

# Factors Influencing Salaries of Data Science Roles

---

## Data Exploration

### Dataset Overview

```
import pandas as pd

salaries_data_frame = pd.read_csv("data/jobs_salaries_2023.csv")
print(salaries_data_frame.shape)

column_types = salaries_data_frame.dtypes
print(column_types)
```

```
(1500, 11)
work_year          int64
experience_level    object
employment_type     object
job_title           object
salary             float64
salary_currency     object
salary_in_usd       float64
employee_residence  object
remote_ratio        int64
company_location    object
company_size        object
dtype: object
```

### Work Year

```
import matplotlib.pyplot as plt

# Count occurrences of each work_year and sort by year
work_year_counts =
salaries_data_frame["work_year"].value_counts().sort_index()


# Calculate percentages
work_year_percent = (work_year_counts / work_year_counts.sum() *
100).round(2)

# Combine into one DataFrame for display
result = pd.DataFrame({"count": work_year_counts, "percentage":
work_year_percent})
print(result)
```

```
# --- Pie Chart ---
plt.figure(figsize=(7, 7))
plt.pie(
    work_year_counts,
    labels=work_year_counts.index,
    autopct="%1.1f%%",
    startangle=90
)
plt.title("Distribution of Records by Work Year")
plt.show()
```

	count	percentage
work_year		
2020	75	5.0
2021	219	14.6
2022	714	47.6
2023	492	32.8

Distribution of Records by Work Year



## Employment Type

```
import matplotlib.pyplot as plt

# Define custom order
order = ["FL", "CT", "PT", "FT"]

# Count occurrences of each employment_type
employment_type_counts =
salaries_data_frame["employment_type"].value_counts()

# Reorder according to the custom order
employment_type_counts = employment_type_counts.reindex(order)

# Calculate percentages
employment_type_percent = (
    employment_type_counts / employment_type_counts.sum() * 100
).round(2)


# Combine into one DataFrame
result = pd.DataFrame(
    {"count": employment_type_counts, "percentage":
employment_type_percent}
)

print(result)

# --- Pie Chart ---
plt.figure(figsize=(7, 7))
plt.pie(
    employment_type_counts,
    labels=employment_type_counts.index,
    autopct="%1.1f%%",
    startangle=90,
    colors=["#66c2a5", "#fc8d62", "#8da0cb", "#e78ac3"],
)
plt.title("Distribution of Employment Types")
plt.show()
```

employment_type	count	percentage
FL	6	0.4
CT	9	0.6
PT	12	0.8
FT	1473	98.2

## Distribution of Employment Types



## Salary

```
import matplotlib.pyplot as plt
import seaborn as sns

salary_in_usd_series = salaries_data_frame["salary_in_usd"]

# --- 1. Descriptive statistics ---
mean_salary = salary_in_usd_series.mean()
median_salary = salary_in_usd_series.median()
min_salary = salary_in_usd_series.min()
max_salary = salary_in_usd_series.max()



print(f"Salary ranges from ${min_salary:,.0f} to ${max_salary:,.0f}")
print(f"Mean salary: ${mean_salary:,.0f}")
print(f"Median salary: ${median_salary:,.0f}")
```

```
# Skew check
skewness = salary_in_usd_series.skew()
print(f"Skewness: {skewness:.2f}")

# --- 2. Histogram ---
plt.figure(figsize=(8, 5))
sns.histplot(salary_in_usd_series, bins=30, kde=True)
plt.axvline(
    mean_salary, color="red", linestyle="--", label=f"Mean =
{mean_salary:,.0f}"
)
plt.axvline(
    median_salary, color="blue", linestyle="--", label=f"Median =
{median_salary:,.0f}"
)
plt.title("Salary Distribution (USD)")
plt.xlabel("Salary")
plt.ylabel("Frequency")
plt.legend()
plt.show()

# --- 3. Boxplot (to reveal outliers) ---
plt.figure(figsize=(6, 3))
sns.boxplot(x=salary_in_usd_series)
plt.title("Salary in USD - Boxplot")
plt.xlabel("Salary in USD")
plt.show()
```

Salary ranges from \$5,409 to \$450,000  
Mean salary: \$130,934  
Median salary: \$130,000  
Skewness: 0.59



Experience Level (with Salary)

```
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# =====
# Preprocessing
# =====
```

```

experience_level_order = ["EN", "MI", "SE", "EX"]
salaries_data_frame["experience_level"] = pd.Categorical(
    salaries_data_frame["experience_level"],
    categories=experience_level_order,
    ordered=True,
)

# =====
# Counts and Percentages
# =====
experience_level_counts = (
    salaries_data_frame["experience_level"].value_counts().sort_index()
)
experience_level_percent = (
    experience_level_counts / experience_level_counts.sum() * 100
).round(1)

result = pd.DataFrame(
    {"count": experience_level_counts, "percentage":
experience_level_percent}
)

# =====
# Mean and Median Salaries
# =====
summary_table = (
    salaries_data_frame.groupby("experience_level")["salary_in_usd"]
    .agg(Mean="mean", Median="median")
    .round(0)
)

# =====
# Combine Tables
# =====
combined_table = result.join(summary_table)
print("Experience Level Summary:")
print(combined_table)

# =====
# Pie Chart
# =====
plt.figure(figsize=(7, 7))
plt.pie(
    experience_level_counts,
    labels=experience_level_counts.index,
    autopct="%1.1f%%",
    startangle=90,
    colors=["#66c2a5", "#fc8d62", "#8da0cb", "#e78ac3"],
)
plt.title("Distribution of Experience Levels")
plt.show()

# =====
# Boxplot

```

```
# =====
plt.figure(figsize=(8, 6))
sns.boxplot(
    data=salaries_data_frame,
    x="experience_level",
    y="salary_in_usd",
    order=experience_level_order,
    palette="Set2",
    hue="experience_level",
    legend=False,
)
plt.title("Salary Distribution by Experience Level", fontsize=14)
plt.ylabel("Salary (USD)")
plt.xlabel("Experience Level")
plt.tight_layout()
plt.show()
```


#### Experience Level Summary:


	count	percentage	Mean	Median
experience_level				
EN	167	11.1	69627.0	60000.0
MI	353	23.5	95473.0	84053.0
SE	922	61.5	151640.0	145000.0
EX	58	3.9	192463.0	188518.0

/var/folders/jh/z981c7zj0vz0gmyfc8mhdxd0000gn/T/ipykernel\_11047/3124317989.py:33: FutureWarning: The default of observed=False is deprecated and will be changed to True in a future version of pandas. Pass observed=False to retain current behavior or observed=True to adopt the future default and silence this warning.

```
salaries_data_frame.groupby("experience_level")["salary_in_usd"]
```







### Job Title (with Salary)

```
import pandas as pd

# =====
# Job title counts & percentages
# =====
job_counts = salaries_data_frame["job_title"].value_counts()
total_jobs = job_counts.sum()

# DataFrame with count and percentage
all_job_titles = job_counts.to_frame("count").reset_index()
all_job_titles.columns = ["job_title", "count"]
all_job_titles["percentage"] = (all_job_titles["count"] / total_jobs *
100).round(2)

# Add accumulated count and percentage
all_job_titles["accumulated_count"] = all_job_titles["count"].cumsum()
all_job_titles["accumulated_percentage"] = (
    all_job_titles["accumulated_count"] / total_jobs * 100
).round(2)

# =====
# Mean and median salary
```

```
# =====
salary_summary = (
    salaries_data_frame.groupby("job_title")["salary_in_usd"]
    .agg(Mean="mean", Median="median")
    .round(0)
    .reset_index()
)

# =====
# Combine tables
# =====
all_job_titles = all_job_titles.merge(salary_summary, on="job_title",
how="left")

# Print summary
print(f"Number of unique job titles: {len(all_job_titles)}")
print("\nAll job titles with accumulation, mean & median salary:")
print(all_job_titles.to_string(index=False, line_width=10000))
```

Number of unique job titles: 69

All job titles with accumulation, mean & median salary:

	job_title	count	percentage	Mean	Median
	Data Engineer	366			24.40
366		24.40	131523.0	130000.0	
	Data Scientist	360			24.00
726		48.40	135659.0	141300.0	
	Data Analyst	210			14.00
936		62.40	100195.0	100000.0	
	Machine Learning Engineer	106			7.07
1042		69.47	145421.0	141942.0	
	Analytics Engineer	53			3.53
1095		73.00	159451.0	152700.0	
	Data Architect	39			2.60
1134		75.60	165886.0	167500.0	
	Applied Scientist	30			2.00
1164		77.60	189030.0	184000.0	
	Data Science Manager	29			1.93
1193		79.53	177154.0	175100.0	
	Research Scientist	27			1.80
1220		81.33	127143.0	102772.0	
	Machine Learning Scientist	17			1.13
1237		82.47	164900.0	180000.0	
	Research Engineer	14			0.93
1251		83.40	184365.0	179500.0	
	Computer Vision Engineer	12			0.80
1263		84.20	139076.0	147500.0	
	ML Engineer	12			0.80
1275		85.00	114463.0	80682.0	

	Data Analytics Manager	11	0.73
1286	85.73 133176.0 140000.0		
	BI Data Analyst	10	0.67
1296	86.40 56665.0 51900.0		
	AI Scientist	10	0.67
1306	87.07 89447.0 50448.0		
	Director of Data Science	10	0.67
1316	87.73 202086.0 180018.0		
	Business Data Analyst	10	0.67
1326	88.40 80750.0 84566.0		
	Applied Machine Learning Scientist	9	0.60
1335	89.00 114501.0 75000.0		
	Big Data Engineer	9	0.60
1344	89.60 51565.0 48289.0		
	ETL Developer	8	0.53
1352	90.13 125192.0 93635.0		
	Data Manager	8	0.53
1360	90.67 124000.0 117500.0		
	Principal Data Scientist	8	0.53
1368	91.20 198171.0 164630.0		
	Applied Data Scientist	8	0.53
1376	91.73 127158.0 89178.0		
	Head of Data	7	0.47
1383	92.20 199780.0 230000.0		
	Data Science Consultant	7	0.47
1390	92.67 69421.0 76833.0		
	Data Specialist	7	0.47
1397	93.13 130000.0 130000.0		
	Data Operations Engineer	6	0.40
1403	93.53 80000.0 80000.0		
	AI Developer	6	0.40
1409	93.93 169670.0 154000.0		
	Lead Data Engineer	6	0.40
1415	94.33 139230.0 120111.0		
	Lead Data Scientist	5	0.33
1420	94.67 87416.0 61566.0		
	Machine Learning Infrastructure Engineer	5	0.33
1425	95.00 127133.0 148800.0		
	Machine Learning Developer	5	0.33
1430	95.33 89726.0 76814.0		
	Data Quality Analyst	5	0.33
1435	95.67 92000.0 100000.0		
	Business Intelligence Engineer	4	0.27
1439	95.93 174150.0 171150.0		
	Computer Vision Software Engineer	4	0.27
1443	96.20 83705.0 82873.0		
	Head of Data Science	4	0.27
1447	96.47 146719.0 138938.0		
	Data Analytics Engineer	4	0.27
1451	96.73 64799.0 64598.0		
	Lead Data Analyst	4	0.27

1455	97.00	86152.0	77500.0		
	Product Data Analyst	3	0.20		
1458	97.20	55357.0	20000.0		
	Data Science Engineer	3	0.20		
1461	97.40	75803.0	60000.0		
	Principal Data Engineer	2	0.13		
1463	97.53	192500.0	192500.0		
	Lead Machine Learning Engineer	2	0.13		
1465	97.67	89720.0	89720.0		
	Cloud Data Engineer	2	0.13		
1467	97.80	124647.0	124647.0		
	Principal Data Analyst	2	0.13		
1469	97.93	122500.0	122500.0		
	ETL Engineer	2	0.13		
1471	98.07	71394.0	71394.0		
	Data Operations Analyst	2	0.13		
1473	98.20	73500.0	73500.0		
	Financial Data Analyst	2	0.13		
1475	98.33	87500.0	87500.0		
	Data Modeler	2	0.13		
1477	98.47	118900.0	118900.0		
	Machine Learning Research Engineer	2	0.13		
1479	98.60	16086.0	16086.0		
	Data Strategist	2	0.13		
1481	98.73	81000.0	81000.0		
	MLOps Engineer	2	0.13		
1483	98.87	129000.0	129000.0		
	Data DevOps Engineer	1	0.07		
1484	98.93	53654.0	53654.0		
	BI Data Engineer	1	0.07		
1485	99.00	60000.0	60000.0		
	Staff Data Scientist	1	0.07		
1486	99.07	105000.0	105000.0		
	Big Data Architect	1	0.07		
1487	99.13	99703.0	99703.0		
	Staff Data Analyst	1	0.07		
1488	99.20	15000.0	15000.0		
	Marketing Data Analyst	1	0.07		
1489	99.27	88654.0	88654.0		
	3D Computer Vision Researcher	1	0.07		
1490	99.33	5409.0	5409.0		
	Machine Learning Researcher	1	0.07		
1491	99.40	50000.0	50000.0		
	Machine Learning Manager	1	0.07		
1492	99.47	117104.0	117104.0		
	Applied Machine Learning Engineer	1	0.07		
1493	99.53	69751.0	69751.0		
	Data Analytics Lead	1	0.07		
1494	99.60	405000.0	405000.0		
	Compliance Data Analyst	1	0.07		
1495	99.67	30000.0	30000.0		

	Data Analytics Consultant	1	0.07
1496	99.73 113000.0 113000.0		
	Head of Machine Learning	1	0.07
1497	99.80 76309.0 76309.0		
	NLP Engineer	1	0.07
1498	99.87 60000.0 60000.0		
	Cloud Data Architect	1	0.07
1499	99.93 250000.0 250000.0		
	Finance Data Analyst	1	0.07
1500	100.00 61896.0 61896.0		

```
import matplotlib.pyplot as plt
import pandas as pd


# Count all job titles and calculate percentages
job_counts = salaries_data_frame["job_title"].value_counts()
total_jobs = job_counts.sum()

# Keep top 10 job titles and group the rest as 'Other'
top_n = 10
top_jobs = job_counts.head(top_n)
other_count = total_jobs - top_jobs.sum()

# Use pd.concat to combine top jobs and 'Other'
job_counts_for_pie = pd.concat([top_jobs, pd.Series({"Other":
other_count})])

# Pie Chart
plt.figure(figsize=(8, 8))
plt.pie(
    job_counts_for_pie,
    labels=job_counts_for_pie.index,
    autopct="%1.1f%%",
    startangle=90,
    colors=plt.cm.tab20.colors # color map for slices
)
plt.title(f"Top {top_n} Job Titles Distribution")
plt.show()
```

## Top 10 Job Titles Distribution



## Company Location

```
import pandas as pd
import matplotlib.pyplot as plt

# =====
# Count company locations
# =====
location_counts = salaries_data_frame["company_location"].value_counts()
total_locations = location_counts.sum()

# Create DataFrame with count and percentage
all_locations = location_counts.to_frame("count").reset_index()
all_locations.columns = ["company_location", "count"]
all_locations["percentage"] = (all_locations["count"] / total_locations *
100).round(2)

# Add accumulated count and percentage
all_locations["accumulated_count"] = all_locations["count"].cumsum()
all_locations["accumulated_percentage"] = (
```

```

    all_locations["accumulated_count"] / total_locations * 100
).round(2)

# =====
# Mean and Median Salaries
# =====
salary_summary = (
    salaries_data_frame.groupby("company_location")["salary_in_usd"]
    .agg(Mean="mean", Median="median")
    .round(0)
    .reset_index()
)

# Merge with main table
all_locations = all_locations.merge(salary_summary, on="company_location",
how="left")

# =====
# Print summary
# =====
print(f"Number of unique company locations: {len(all_locations)}")

pd.set_option("display.max_rows", None) # Show all rows in output
print("\nAll company locations with salary stats:")
print(all_locations.to_string(index=False, line_width=10000))

# =====
# Pie Chart
# =====
plt.figure(figsize=(8, 8))
plt.pie(
    all_locations["count"],
    labels=all_locations["company_location"], # now from column
    autopct="%1.1f%%",
    startangle=90,
    counterclock=False,
)
plt.title("Company Location Distribution", fontsize=14)
plt.show()

```

Number of unique company locations: 58

All company locations with salary stats:

company_location	count	percentage	accumulated_count
US	1101	73.40	1101
GB	91	6.07	1192
CA	50	3.33	1242




		IN	37	2.47	1279
85.27	33720.0	20670.0			
		DE	35	2.33	1314
87.60	86249.0	76833.0			
		ES	29	1.93	1343
89.53	50044.0	47282.0			
		FR	20	1.33	1363
90.87	61112.0	55196.0			
		GR	13	0.87	1376
91.73	54786.0	52533.0			
		NL	11	0.73	1387
92.47	71873.0	69741.0			
		JP	6	0.40	1393
92.87	114127.0	75682.0			
		PT	6	0.40	1399
93.27	40065.0	40062.0			
		BR	5	0.33	1404
93.60	13975.0	12901.0			
		AT	5	0.33	1409
93.93	67765.0	61989.0			
		AU	5	0.33	1414
94.27	100834.0	83864.0			
		PL	4	0.27	1418
94.53	65587.0	40103.0			
		BE	4	0.27	1422
94.80	76865.0	83398.0			
		DK	4	0.27	1426
95.07	45558.0	37252.0			
		TR	4	0.27	1430
95.33	21322.0	22586.0			
		PR	4	0.27	1434
95.60	167500.0	167500.0			
		NG	4	0.27	1438
95.87	47500.0	40000.0			
		UA	4	0.27	1442
96.13	57850.0	55000.0			
		IE	3	0.20	1445
96.33	117764.0	113750.0			
		PK	3	0.20	1448
96.53	13333.0	12000.0			
		FI	3	0.20	1451
96.73	68793.0	68318.0			
		LU	3	0.20	1454
96.93	43943.0	59102.0			
		AE	3	0.20	1457
97.13	100000.0	115000.0			
		CH	3	0.20	1460
97.33	60940.0	56536.0			
		IT	2	0.13	1462
97.47	36366.0	36366.0			
		RU	2	0.13	1464

97.60	157500.0	157500.0		
	SI	2	0.13	1466
97.73	63831.0	63831.0		
	MX	2	0.13	1468
97.87	46756.0	46756.0		
	CF	2	0.13	1470
98.00	48609.0	48609.0		
	CZ	2	0.13	1472
98.13	50234.0	50234.0		
	SG	2	0.13	1474
98.27	77276.0	77276.0		
	ID	2	0.13	1476
98.40	34208.0	34208.0		
	AS	2	0.13	1478
98.53	34026.0	34026.0		
	CO	1	0.07	1479
98.60	21844.0	21844.0		
	HU	1	0.07	1480
98.67	35735.0	35735.0		
	KE	1	0.07	1481
98.73	9272.0	9272.0		
	TH	1	0.07	1482
98.80	15000.0	15000.0		
	NZ	1	0.07	1483
98.87	125000.0	125000.0		
	CL	1	0.07	1484
98.93	40038.0	40038.0		
	MD	1	0.07	1485
99.00	18000.0	18000.0		
	HR	1	0.07	1486
99.07	45618.0	45618.0		
	IL	1	0.07	1487
99.13	119059.0	119059.0		
	CN	1	0.07	1488
99.20	100000.0	100000.0		
	EE	1	0.07	1489
99.27	31520.0	31520.0		
	IQ	1	0.07	1490
99.33	100000.0	100000.0		
	RO	1	0.07	1491
99.40	60000.0	60000.0		
	DZ	1	0.07	1492
99.47	100000.0	100000.0		
	HN	1	0.07	1493
99.53	20000.0	20000.0		
	HK	1	0.07	1494
99.60	65062.0	65062.0		
	MY	1	0.07	1495
99.67	40000.0	40000.0		
	EG	1	0.07	1496
99.73	22800.0	22800.0		

	AR	1	0.07	1497
99.80	50000.0	50000.0		
	PH	1	0.07	1498
99.87	50000.0	50000.0		
	B0	1	0.07	1499
99.93	7500.0	7500.0		
	MT	1	0.07	1500
100.00	28369.0	28369.0		

Company Location Distribution



Employee Residence

```
import pandas as pd
import matplotlib.pyplot as plt

# =====
```

```

# Count company locations
# =====
location_counts = salaries_data_frame["employee_residence"].value_counts()
total_locations = location_counts.sum()

# Create DataFrame with count and percentage
all_locations = location_counts.to_frame("count").reset_index()
all_locations.columns = ["employee_residence", "count"]
all_locations["percentage"] = (all_locations["count"] / total_locations *
100).round(2)

# Add accumulated count and percentage
all_locations["accumulated_count"] = all_locations["count"].cumsum()
all_locations["accumulated_percentage"] = (
    all_locations["accumulated_count"] / total_locations * 100
).round(2)

# =====
# Mean and Median Salaries
# =====
salary_summary = (
    salaries_data_frame.groupby("employee_residence")["salary_in_usd"]
    .agg(Mean="mean", Median="median")
    .round(0)
    .reset_index()
)

# Merge with main table
all_locations = all_locations.merge(salary_summary,
on="employee_residence", how="left")

# =====
# Print summary
# =====
print(f"Number of unique company locations: {len(all_locations)}")

pd.set_option("display.max_rows", None) # Show all rows in output
print("\nAll company locations with salary stats:")
print(all_locations.to_string(index=False, line_width=10000))

# =====
# Pie Chart
# =====
plt.figure(figsize=(8, 8))
plt.pie(
    all_locations["count"],
    labels=all_locations["employee_residence"], # now from column
    autopct="%1.1f%%",
    startangle=90,
    counterclock=False,
)
plt.title("Company Location Distribution", fontsize=14)
plt.show()

```

Number of unique company locations: 61


All company locations with salary stats:

employee_residence	count	percentage	accumulated_count	accumulated_percentage	Mean	Median
US	1074	71.60	1074			
GB	86	5.73	1160			
CA	48	3.20	1208			
IN	46	3.07	1254			
ES	31	2.07	1285			
DE	30	2.00	1315			
FR	24	1.60	1339			
GR	15	1.00	1354			
NL	12	0.80	1366			
PT	10	0.67	1376			
BR	8	0.53	1384			
JP	7	0.47	1391			
PK	6	0.40	1397			
AU	6	0.40	1403			
NG	5	0.33	1408			
PR	5	0.33	1413			
PL	4	0.27	1417			
IT	4	0.27	1421			
TR	4	0.27	1425			
AT	4	0.27	1429			
BE	4	0.27	1433			
RU	4	0.27	1437			

		UA	4	0.27	1441
96.07	57850.0	55000.0			
		B0	3	0.20	1444
96.27	52500.0	75000.0			
		DK	3	0.20	1447
96.47	31193.0	28609.0			
		AR	3	0.20	1450
96.67	52667.0	50000.0			
		IE	3	0.20	1453
96.87	117764.0	113750.0			
		SG	3	0.20	1456
97.07	91203.0	89294.0			
		AE	3	0.20	1459
97.27	100000.0	115000.0			
		SI	2	0.13	1461
97.40	63831.0	63831.0			
		CH	2	0.13	1463
97.53	88469.0	88469.0			
		CF	2	0.13	1465
97.67	48609.0	48609.0			
		R0	2	0.13	1467
97.80	51419.0	51419.0			
		HK	2	0.13	1469
97.93	65542.0	65542.0			
		VN	2	0.13	1471
98.07	44200.0	44200.0			
		FI	2	0.13	1473
98.20	69030.0	69030.0			
		PH	2	0.13	1475
98.33	47880.0	47880.0			
		HU	2	0.13	1477
98.47	35997.0	35997.0			
		RS	1	0.07	1478
98.53	25532.0	25532.0			
		JE	1	0.07	1479
98.60	100000.0	100000.0			
		KE	1	0.07	1480
98.67	9272.0	9272.0			
		LU	1	0.07	1481
98.73	59102.0	59102.0			
		C0	1	0.07	1482
98.80	21844.0	21844.0			
		NZ	1	0.07	1483
98.87	125000.0	125000.0			
		CL	1	0.07	1484
98.93	40038.0	40038.0			
		MD	1	0.07	1485
99.00	18000.0	18000.0			
		HR	1	0.07	1486
99.07	45618.0	45618.0			
		MX	1	0.07	1487

99.13	33511.0	33511.0			
		EG	1	0.07	1488
99.20	22800.0	22800.0			
		BG	1	0.07	1489
99.27	80000.0	80000.0			
		IQ	1	0.07	1490
99.33	100000.0	100000.0			
		DZ	1	0.07	1491
99.40	100000.0	100000.0			
		CZ	1	0.07	1492
99.47	69999.0	69999.0			
		TN	1	0.07	1493
99.53	30469.0	30469.0			
		HN	1	0.07	1494
99.60	20000.0	20000.0			
		EE	1	0.07	1495
99.67	31520.0	31520.0			
		MY	1	0.07	1496
99.73	200000.0	200000.0			
		ID	1	0.07	1497
99.80	15000.0	15000.0			
		DO	1	0.07	1498
99.87	110000.0	110000.0			
		TH	1	0.07	1499
99.93	15000.0	15000.0			
		MT	1	0.07	1500
100.00	28369.0	28369.0			

## Company Location Distribution



## Remote Ratio (with Salary)

```
import pandas as pd
import matplotlib.pyplot as plt

# =====
# Define custom order
# =====
order = [0, 50, 100]

# Count occurrences of each remote_ratio
remote_ratio_counts = salaries_data_frame["remote_ratio"].value_counts()

# Reorder according to the custom order
remote_ratio_counts = remote_ratio_counts.reindex(order)
```



```

# Calculate percentages
remote_ratio_percent = (remote_ratio_counts / remote_ratio_counts.sum() *
100).round(2)

# =====
# Mean and Median Salaries
# =====
salary_summary = (
    salaries_data_frame.groupby("remote_ratio")["salary_in_usd"]
    .agg(Mean="mean", Median="median")
    .round(0)
    .reindex(order)
)

# =====
# Combine into one DataFrame
# =====
result = pd.DataFrame(
    {
        "count": remote_ratio_counts,
        "percentage": remote_ratio_percent,
        "Mean": salary_summary["Mean"],
        "Median": salary_summary["Median"],
    }
)

print("Remote Work Ratio Summary:")
print(result)


# =====
# Draw Pie Chart
# =====
plt.figure(figsize=(6, 6))
plt.pie(
    result["count"],
    labels=result.index.astype(str) + " % remote",
    autopct="%1.1f%%",
    startangle=90,
    counterclock=False,
)
plt.title("Remote Work Ratio Distribution", fontsize=14)
plt.show()

```

Remote Work Ratio Summary:

	count	percentage	Mean	Median
remote_ratio				
0	579	38.60	143867.0	139430.0
50	130	8.67	81360.0	65135.0
100	791	52.73	129658.0	131050.0

## Remote Work Ratio Distribution



## Company Size (with Salary)

```
import pandas as pd
import matplotlib.pyplot as plt

# =====
# Define custom order
# =====
order = ["S", "M", "L"]

# Count occurrences of each company_size
company_size_counts = salaries_data_frame["company_size"].value_counts()

# Reorder according to the custom order
company_size_counts = company_size_counts.reindex(order)

# Calculate percentages
company_size_percent = (company_size_counts / company_size_counts.sum() *
100).round(2)

# =====
# Mean and Median Salaries
# =====
salary_summary = (
```

```

    salaries_data_frame.groupby("company_size")["salary_in_usd"]
    .agg(Mean="mean", Median="median")
    .round(0)
    .reindex(order)
)

# =====
# Combine into one DataFrame
# =====
result = pd.DataFrame(
    {
        "count": company_size_counts,
        "percentage": company_size_percent,
        "Mean": salary_summary["Mean"],
        "Median": salary_summary["Median"],
    }
)

print("Company Size Summary:")
print(result)


# =====
# Pie chart visualization
# =====
plt.figure(figsize=(6, 6))
plt.pie(
    company_size_counts,
    labels=order,
    autopct="%1.1f%%",
    startangle=90,
    counterclock=False,
)
plt.title("Distribution of Company Size")
plt.show()

```

#### Company Size Summary:

	count	percentage	Mean	Median
company_size				
S	107	7.13	77723.0	61566.0
M	1073	71.53	139114.0	137270.0
L	320	21.33	121396.0	112300.0

Distribution of Company Size



## Data Preparation and Model Training

### Handling Data Issues

#### Missing Values Issues

```
# Drop rows with any missing values
salaries_data_frame = salaries_data_frame.dropna()
salaries_data_frame.shape
```

```
(1494, 11)
```

#### Employment Type Filter

```
# Keep only full-time employees
salaries_data_frame =
salaries_data_frame[salaries_data_frame["employment_type"] == "FT"] #
assuming "FT" is the code for full-time
salaries_data_frame.shape
```

```
(1467, 11)
```

## Job Titles Filter

```
# ===== FILTER OUT JOB TITLES WITH FEWER THAN 100 RECORDS =====  
threshold = 100  
job_counts = salaries_data_frame["job_title"].value_counts()  
salaries_data_frame =  
salaries_data_frame[salaries_data_frame["job_title"].isin(job_counts[job_c  
ounts >= threshold].index)]  
salaries_data_frame.shape
```

```
(1030, 11)
```

## Salary Distribution after Filter

```
import pandas as pd  
import matplotlib.pyplot as plt  
import seaborn as sns  
  
# --- 1. Descriptive statistics ---  
salary_in_usd_series = salaries_data_frame["salary_in_usd"]  
  
mean_salary = salary_in_usd_series.mean()  
median_salary = salary_in_usd_series.median()  
min_salary = salary_in_usd_series.min()  
max_salary = salary_in_usd_series.max()  
  
print(f"Salary ranges from ${min_salary:,.0f} to ${max_salary:,.0f}")  
print(f"Mean salary: ${mean_salary:,.0f}")  
print(f"Median salary: ${median_salary:,.0f}")  
  
# Skew check  
skewness = salary_in_usd_series.skew()  
print(f"Skewness: {skewness:.2f}")  
  
# --- 2. Histogram ---  
plt.figure(figsize=(8, 5))  
sns.histplot(salary_in_usd_series, bins=30, kde=True)  
plt.axvline(  
    mean_salary, color="red", linestyle="--", label=f"Mean =  
{mean_salary:,.0f}"  
)
```

```

plt.axvline(
    median_salary, color="blue", linestyle="--", label=f"Median = {median_salary:,.0f}"
)
plt.title("Salary Distribution (USD)")
plt.xlabel("Salary")
plt.ylabel("Frequency")
plt.legend()
plt.show()

# --- 3. Boxplot (to reveal outliers) ---
plt.figure(figsize=(6, 3))
sns.boxplot(x=salary_in_usd_series)
plt.title("Salary in USD - Boxplot")
plt.xlabel("Salary in USD")
plt.show()

# --- 4. Identify outliers using IQR rule ---
Q1 = salary_in_usd_series.quantile(0.25)
Q3 = salary_in_usd_series.quantile(0.75)
IQR = Q3 - Q1

lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR

outliers = salaries_data_frame[(salaries_data_frame["salary_in_usd"] <
lower_bound) | (salaries_data_frame["salary_in_usd"] > upper_bound)]

print(f"\nNumber of outliers: {len(outliers)}")
print("Outlier rows:")
print(outliers.sort_values("salary_in_usd").to_string(index=False,
line_width=10000))

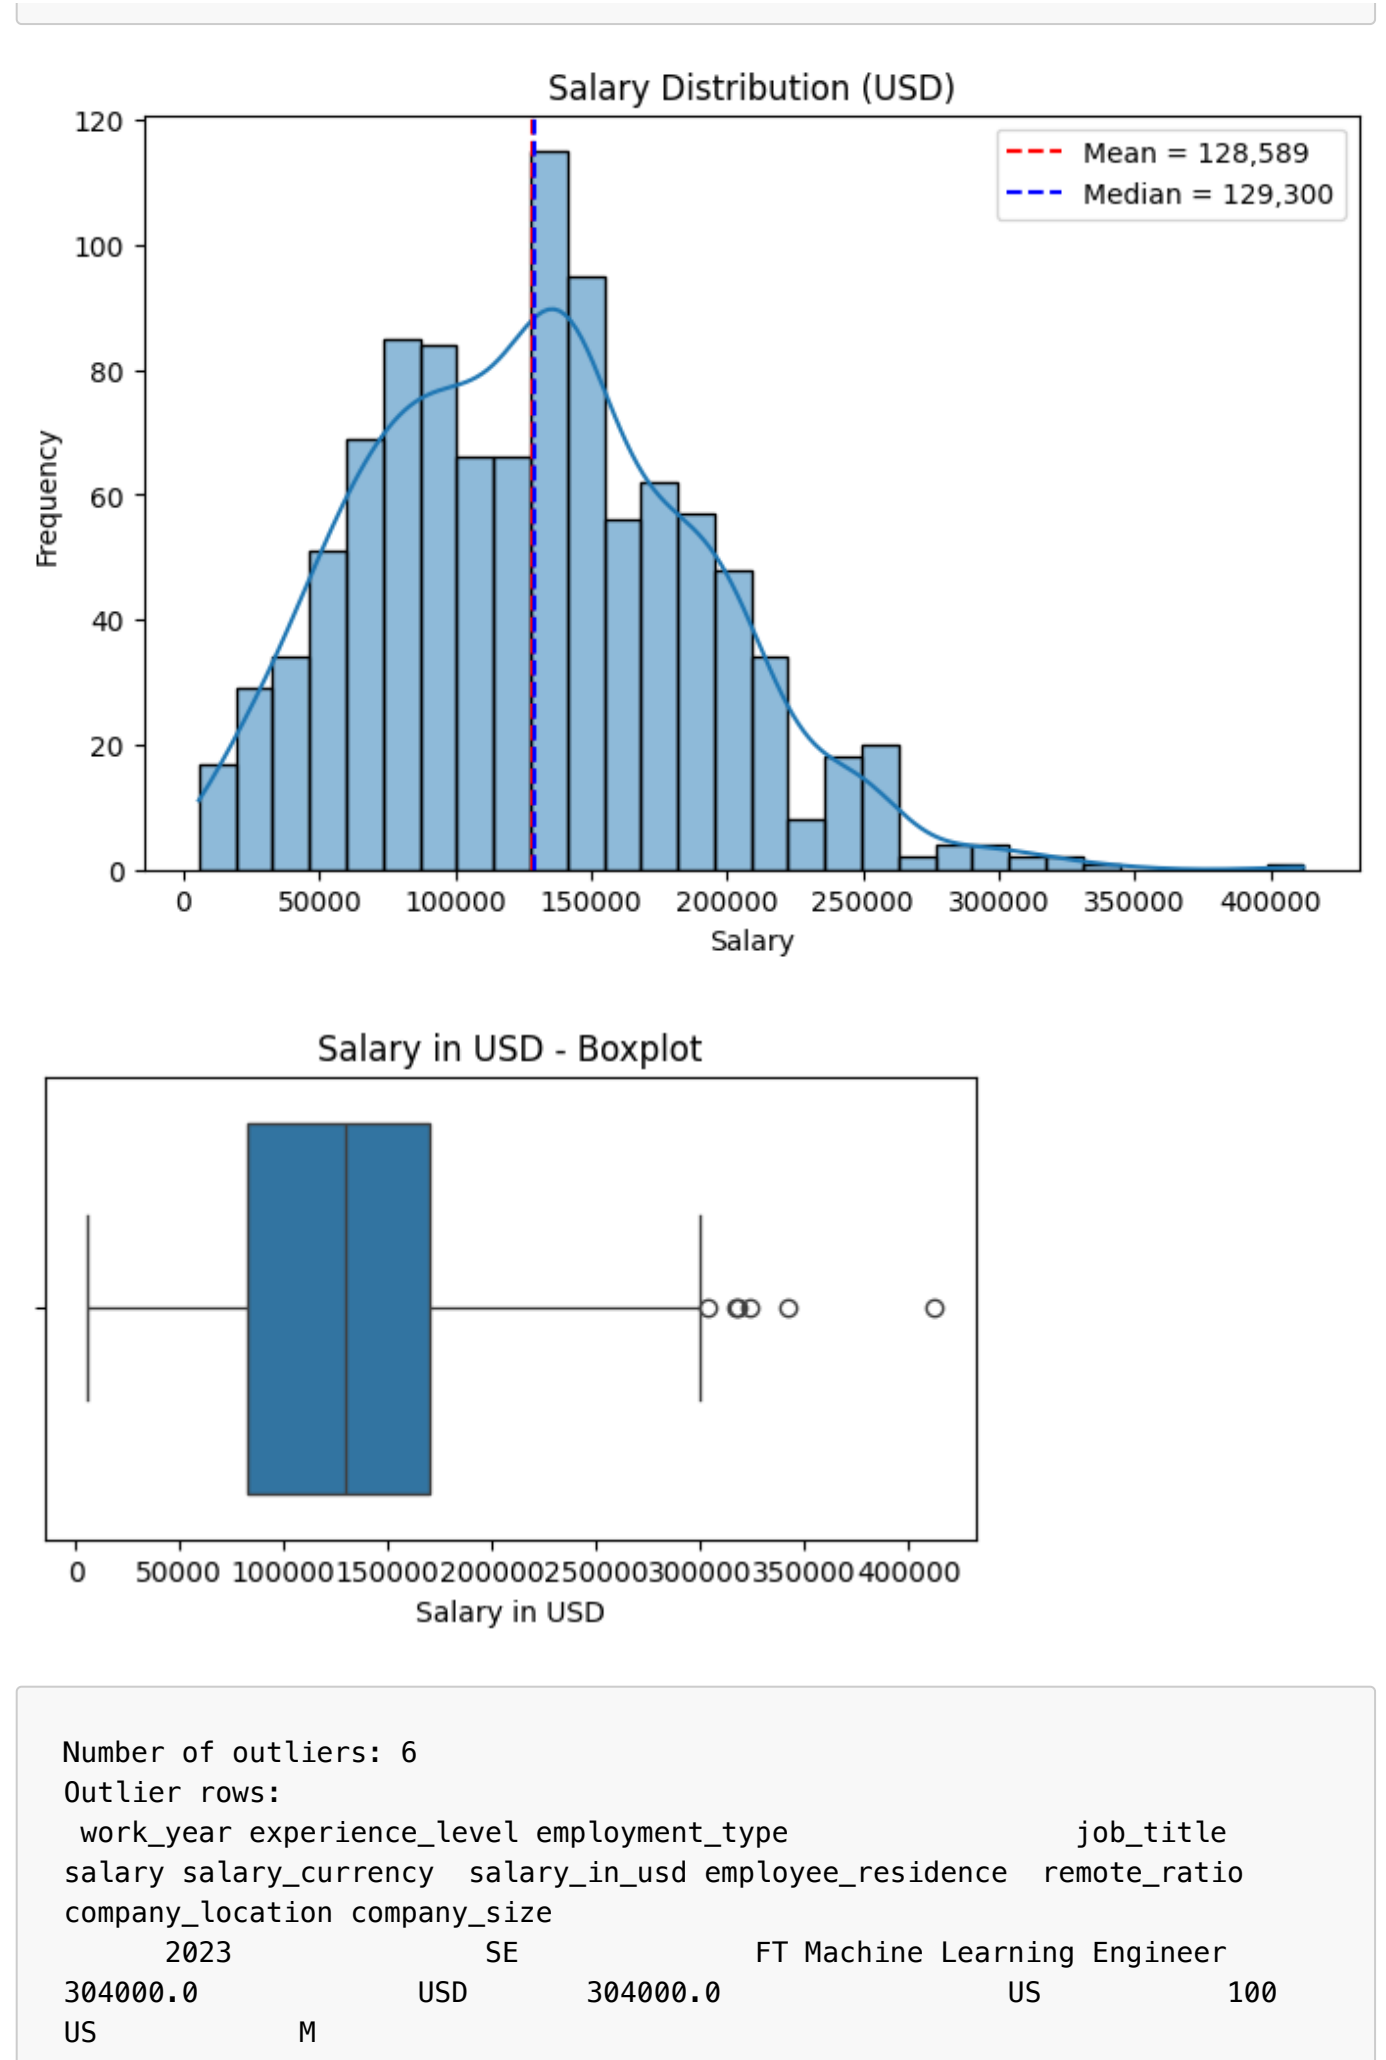
# --- 5. Remove outliers ---
print(f"\nData shape before removing outliers: {salaries_data_frame.shape}")

salaries_data_frame = salaries_data_frame[
    (salaries_data_frame["salary_in_usd"] >= lower_bound) &
    (salaries_data_frame["salary_in_usd"] <= upper_bound)
].copy()

print(f"\nData shape after removing outliers: {salaries_data_frame.shape}")

```

Salary ranges from \$5,679 to \$412,000  
 Mean salary: \$128,589  
 Median salary: \$129,300  
 Skewness: 0.43



2023	SE	FT	Data Scientist
317070.0	USD	317070.0	US 0
US	M		
2023	SE	FT Machine Learning Engineer	
318300.0	USD	318300.0	US 100
US	M		
2022	EX	FT	Data Engineer
324000.0	USD	324000.0	US 100
US	M		
2023	SE	FT Machine Learning Engineer	
342300.0	USD	342300.0	US 0
US	L		
2020	SE	FT	Data Scientist
412000.0	USD	412000.0	US 100
US	L		

Data shape before removing outliers: (1030, 11)

Data shape after removing outliers: (1024, 11)

## Model Training and Performance (Supervised Learning)

### First Try

```
# =====
# 1. Import libraries
# =====
from sklearn.model_selection import train_test_split
from sklearn.compose import ColumnTransformer
import matplotlib.pyplot as plt
import seaborn as sns
from xgboost import XGBRegressor
from lightgbm import LGBMRegressor
from catboost import CatBoostRegressor
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor,
GradientBoostingRegressor
from sklearn.metrics import root_mean_squared_error, r2_score,
mean_absolute_error
from sklearn.preprocessing import StandardScaler, OneHotEncoder

print(salaries_data_frame.shape)

# =====
# 2. Define features & target
# =====
categorical_cols = [
    "experience_level",
    "job_title",
    "employee_residence",
```



```
        "company_location",
        "company_size",
    ]
    numeric_cols = ["remote_ratio"]
    features = categorical_cols + numeric_cols

    X = salaries_data_frame[features]
    X.shape

    y = salaries_data_frame["salary_in_usd"]
    y.shape

    # =====
    # 3. Split dataset
    # =====
    X_train, X_test, y_train, y_test = train_test_split(
        X, y, test_size=0.2, random_state=42
    )

    # =====
    # 4. Preprocess features
    # =====
    preprocessor_1 = ColumnTransformer(
        transformers=[
            ("cat", OneHotEncoder(handle_unknown="ignore"), categorical_cols),
            ("num", StandardScaler(), numeric_cols),
        ]
    )

    X_train_processed = preprocessor_1.fit_transform(X_train)
    X_test_processed = preprocessor_1.transform(X_test)

    # =====
    # 5. Define models_1
    # =====
    models_1 = {
        "LinearRegression": LinearRegression(),
        "RandomForest": RandomForestRegressor(
            n_estimators=500,
            max_depth=10,
            min_samples_leaf=2,
            max_features="sqrt",
            random_state=42,
            n_jobs=-1,
        ),
        "GradientBoosting": GradientBoostingRegressor(
            n_estimators=500, learning_rate=0.05, max_depth=5, random_state=42
        ),
        "XGBoost": XGBRegressor(
            n_estimators=500,
            learning_rate=0.05,
            max_depth=6,
            subsample=0.8,
            colsample_bytree=0.8,
```

```

        random_state=42,
    ),
    "LightGBM": LGBMRegressor(
        n_estimators=500,
        learning_rate=0.05,
        max_depth=-1,
        num_leaves=31,
        subsample=0.8,
        colsample_bytree=0.8,
        random_state=42,
        verbose=-1,
    ),
    "CatBoost": CatBoostRegressor(
        iterations=500, learning_rate=0.05, depth=6, random_state=42,
        verbose=0
    ),
}
# =====
# 6. Train & evaluate
# =====
results = {}
true_avg_salary = y_test.mean() # True average salary

for name, model in models_1.items():
    model.fit(X_train_processed, y_train)
    y_pred = model.predict(X_test_processed)

    rmse = root_mean_squared_error(y_test, y_pred)
    mae = mean_absolute_error(y_test, y_pred)
    r2 = r2_score(y_test, y_pred)
    pred_avg_salary = y_pred.mean() # Predicted average salary

    results[name] = {
        "RMSE": rmse,
        "RMSE % of Avg": (rmse / true_avg_salary) * 100,
        "MAE": mae,
        "MAE % of Avg": (mae / true_avg_salary) * 100,
        "R²": r2,
        "True Avg Salary": true_avg_salary,
        "Predicted Avg Salary": pred_avg_salary,
    }

# Convert to DataFrame for easy comparison
results_df = pd.DataFrame(results).T
results_df = results_df.sort_values(by="R²", ascending=False)

print(results_df.round(2).to_string(line_width=10000))
print("\nBest model based on R²:", results_df.index[0])

```

(1024, 11)

	RMSE	RMSE % of Avg	MAE	MAE % of Avg	R²
--	------	---------------	-----	--------------	----

Model	True Avg Salary	Predicted Avg Salary	RMSE	R <sup>2</sup>
LinearRegression	35762.74	29.00 27385.69	22.21	0.62
123309.06	120824.26			
RandomForest	36751.52	29.80 29248.37	23.72	0.60
123309.06	124106.12			
CatBoost	36966.07	29.98 28581.21	23.18	0.59
123309.06	124048.19			
LightGBM	37853.95	30.70 30307.25	24.58	0.58
123309.06	124148.50			
GradientBoosting	37888.40	30.73 28769.98	23.33	0.57
123309.06	124405.32			
XGBoost	39276.33	31.85 29720.77	24.10	0.54
123309.06	125220.46			

Best model based on R<sup>2</sup>: LinearRegression

```
/opt/homebrew/lib/python3.13/site-
packages/sklearn/utils/validation.py:2739: UserWarning: X does not have
valid feature names, but LGBMRegressor was fitted with feature names
warnings.warn(
```

## Linear Regression

```
import matplotlib.pyplot as plt

# Fit Linear Regression only
lin_model = LinearRegression()
lin_model.fit(X_train_processed, y_train)

# Predictions
y_pred_lin = lin_model.predict(X_test_processed)

# Evaluation
rmse = root_mean_squared_error(y_test, y_pred_lin)
r2 = r2_score(y_test, y_pred_lin)


print(f"Linear Regression → RMSE: {rmse:.2f}, R²: {r2:.3f}")

# ===== Plot Predicted vs Actual =====
plt.figure(figsize=(6,6))
plt.scatter(y_test, y_pred_lin, alpha=0.5, edgecolor="k")
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()],
         color="red", linewidth=2, label="Perfect Prediction")

plt.xlabel("Actual Salary (USD)")
plt.ylabel("Predicted Salary (USD)")
plt.title(f"Linear Regression: Predicted vs Actual (R² = {r2:.3f})")
plt.legend()
plt.tight_layout()
```

```
plt.show()
```

Linear Regression → RMSE: 35762.74,  $R^2$ : 0.621



## Features Importance

```
# =====
# 7. Aggregate feature importances by original feature (LinearRegression)
# =====

# Get coefficients from LinearRegression
linreg_coefs = models_1["LinearRegression"].coef_

# Use same OHE feature names as before
ohe = preprocessor_1.named_transformers_["cat"]
ohe_features = ohe.get_feature_names_out(categorical_cols)
```

```

all_features = list(ohe_features) + numeric_cols

# Map back to original columns
def map_to_original(feature_name):
    for col in categorical_cols:
        if feature_name.startswith(col + "_"):
            return col
    if feature_name in numeric_cols:
        return feature_name
    return feature_name

original_features = [map_to_original(f) for f in all_features]

# Aggregate absolute coefficients as importance
feature_importance_salaries_data_frame = (
    pd.DataFrame({"feature": original_features, "importance":
abs(linreg_coefs)})
    .groupby("feature")
    .sum()
    .sort_values(by="importance", ascending=False)
    .reset_index()
)

print("\n=== Aggregated Feature Importances (LinearRegression
coefficients) ===")
print(feature_importance_salaries_data_frame)

# =====
# 8. Visualize aggregated feature importances
# =====
plt.figure(figsize=(8, 5))
sns.barplot(x="importance", y="feature",
data=feature_importance_salaries_data_frame)
plt.title("Aggregated Feature Importances (LinearRegression
Coefficients)")
plt.xlabel("Importance (absolute coefficient)")
plt.ylabel("Feature")
plt.tight_layout()
plt.show()


```

```

=== Aggregated Feature Importances (LinearRegression coefficients) ===

```

	feature	importance
0	employee_residence	870374.002466
1	company_location	805913.144172
2	experience_level	116335.128606
3	job_title	66494.832056
4	company_size	22893.529989
5	remote_ratio	42.901199



## Second Training

```
# =====
# 1. Import libraries
# =====
from sklearn.model_selection import train_test_split
from sklearn.compose import ColumnTransformer
from xgboost import XGBRegressor
from lightgbm import LGBMRegressor
from catboost import CatBoostRegressor
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor,
GradientBoostingRegressor
from sklearn.metrics import root_mean_squared_error, r2_score,
mean_absolute_error
from sklearn.preprocessing import StandardScaler, OneHotEncoder

print(salaries_data_frame.shape)

# =====
# 2. Define features & target
# =====
features_2 = ["experience_level", "job_title", "employee_residence",
"company_location"]

X = salaries_data_frame[features_2]
X.shape
```

```
y = salaries_data_frame["salary_in_usd"]
y.shape

# =====
# 3. Split dataset
# =====
X_train_2, X_test_2, y_train_2, y_test_2 = train_test_split(
    X, y, test_size=0.2, random_state=42
)

# =====
# 4. Preprocess features
# =====
preprocessor_2 = ColumnTransformer(
    transformers=[("cat", OneHotEncoder(handle_unknown="ignore"),
features_2)]
)

X_train_processed_2 = preprocessor_2.fit_transform(X_train_2)
X_test_processed_2 = preprocessor_2.transform(X_test_2)

# =====
# 5. Define models_2
# =====
models_2 = {
    "LinearRegression": LinearRegression(),
    "RandomForest": RandomForestRegressor(
        n_estimators=500,
        max_depth=10,
        min_samples_leaf=2,
        max_features="sqrt",
        random_state=42,
        n_jobs=-1,
    ),
    "GradientBoosting": GradientBoostingRegressor(
        n_estimators=500, learning_rate=0.05, max_depth=5, random_state=42
    ),
    "XGBoost": XGBRegressor(
        n_estimators=500,
        learning_rate=0.05,
        max_depth=6,
        subsample=0.8,
        colsample_bytree=0.8,
        random_state=42,
    ),
    "LightGBM": LGBMRegressor(
        n_estimators=500,
        learning_rate=0.05,
        max_depth=-1,
        num_leaves=31,
        subsample=0.8,
        colsample_bytree=0.8,
        random_state=42,
        verbose=-1,
    ),
}
```

```

    ),
    "CatBoost": CatBoostRegressor(
        iterations=500, learning_rate=0.05, depth=6, random_state=42,
        verbose=0
    ),
}

# =====
# 6. Train & evaluate
# =====
results_2 = {}
true_avg_salary = y_test.mean() # True average salary

for name, model in models_2.items():
    model.fit(X_train_processed_2, y_train_2)
    y_pred_2 = model.predict(X_test_processed_2)

    rmse_2 = root_mean_squared_error(y_test_2, y_pred_2)
    mae_2 = mean_absolute_error(y_test_2, y_pred_2)
    r2_2 = r2_score(y_test_2, y_pred_2)
    pred_avg_salary_2 = y_pred_2.mean() # Predicted average salary

    results_2[name] = {
        "RMSE": rmse_2,
        "MAE": mae_2,
        "R²": r2_2,
        "Predicted Avg Salary": pred_avg_salary_2,
    }

# Convert to DataFrame for easy comparison
results_df_2 = pd.DataFrame(results_2).T
results_df_2 = results_df_2.sort_values(by="R²", ascending=False)

print(results_df_2)
print("\nBest model based on R²:", results_df_2.index[0])

```

```

(1024, 11)

```

	RMSE	MAE	R²	Predicted Avg
Salary				
LinearRegression	35460.262713	27612.691817	0.627127	
120270.467793				
CatBoost	36032.043269	28061.036067	0.615005	
122543.885647				
GradientBoosting	36277.145824	27990.598879	0.609749	
123513.312532				
XGBoost	36652.996185	28507.822847	0.601621	
123548.078125				
RandomForest	36753.630485	29243.640004	0.599430	
123197.168178				
LightGBM	37258.456403	29757.696103	0.588351	
122402.832104				



Best model based on  $R^2$ : LinearRegression

```
/opt/homebrew/lib/python3.13/site-
packages/sklearn/utils/validation.py:2739: UserWarning: X does not have
valid feature names, but LGBMRegressor was fitted with feature names
warnings.warn(
```

## Feature Importance

```
# =====
# 7. Aggregate feature importances by original feature (LinearRegression)
# =====

# Get coefficients from LinearRegression
linear_regression_coefs = models_2["LinearRegression"].coef_

# Use same OHE feature names as before
ohe_2 = preprocessor_2.named_transformers_["cat"]
ohe_features_2 = ohe_2.get_feature_names_out(features_2)
all_features_2 = list(ohe_features_2)

# Map back to original columns
def map_to_original(feature_name):
    for col in categorical_cols:
        if feature_name.startswith(col + "_"):
            return col
    if feature_name in numeric_cols:
        return feature_name
    return feature_name


original_features_2 = [map_to_original(f) for f in all_features_2]

# Aggregate absolute coefficients as importance
feature_importance_salaries_data_frame = (
    pd.DataFrame(
        {"feature": original_features_2, "importance":
abs(linear_regression_coefs)}
    )
    .groupby("feature")
    .sum()
    .sort_values(by="importance", ascending=False)
    .reset_index()
)

print("\n=== Aggregated Feature Importances (LinearRegression
coefficients) ===")
print(feature_importance_salaries_data_frame)
```

```
# =====
# 8. Visualize aggregated feature importances
# =====
plt.figure(figsize=(8, 5))
sns.barplot(x="importance", y="feature",
data=feature_importance_salaries_data_frame)
plt.title("Aggregated Feature Importances (LinearRegression
Coefficients)")
plt.xlabel("Importance (absolute coefficient)")
plt.ylabel("Feature")
plt.tight_layout()
plt.show()
```

```
=== Aggregated Feature Importances (LinearRegression coefficients) ===
      feature      importance
0  employee_residence  906657.879079
1   company_location  803909.312904
2   experience_level  125013.224016
3         job_title   62849.875465
```



Group Employees by Job Title, Experience Level, Employee Residence, Company Location

```
import pandas as pd
```

```

# Group by multiple columns
grouped = (
    salaries_data_frame.groupby(
        ["job_title", "experience_level", "employee_residence",
"company_location"]
    )["salary_in_usd"]
    .agg(count="count", mean_salary="mean", median_salary="median")
    .reset_index()
)

# Round salaries
grouped["mean_salary"] = grouped["mean_salary"].round(0)
grouped["median_salary"] = grouped["median_salary"].round(0)

# Add percentage column
total_count = grouped["count"].sum()
grouped["percentage"] = (grouped["count"] / total_count * 100).round(2)

# Remove rows with count = 0 (safety check)
grouped = grouped[grouped["count"] > 0]

# Sort by count (descending)
grouped = grouped.sort_values(by="count", ascending=False)

# Reorder columns
grouped = grouped[
    [
        "job_title",
        "experience_level",
        "employee_residence",
        "company_location",
        "count",
        "percentage",
        "mean_salary",
        "median_salary",
    ]
]

# Show result
print(grouped.to_string(index=False, line_width=10000))

```

		job_title	experience_level	employee_residence		
company_location		count	percentage	mean_salary	median_salary	
		Data Engineer		SE		US
US	216	21.09	154309.0	150000.0		
		Data Scientist		SE		US
US	212	20.70	165500.0	160000.0		
		Data Analyst		SE		US
US	114	11.13	121276.0	115467.0		
		Machine Learning Engineer		SE		US
US	52	5.08	177997.0	183000.0		

		Data Engineer		MI	US
US	35	3.42	117558.0	110000.0	
		Data Analyst		MI	US
US	29	2.83	109606.0	110000.0	
		Data Scientist		MI	US
US	28	2.73	129219.0	130000.0	
		Data Engineer		MI	GB
GB	27	2.64	84430.0	82528.0	
		Data Engineer		EN	US
US	16	1.56	82625.0	82500.0	
		Data Scientist		MI	GB
GB	16	1.56	86411.0	78497.0	
		Data Analyst		EN	US
US	15	1.46	74620.0	72000.0	
		Data Engineer		EX	US
US	13	1.27	206038.0	207000.0	
		Data Scientist		EN	US
US	11	1.07	89355.0	90000.0	
		Data Scientist		SE	CA
CA	10	0.98	169443.0	175500.0	
		Data Analyst		MI	GB
GB	9	0.88	50965.0	49253.0	
		Machine Learning Engineer		EN	US
US	8	0.78	137635.0	131500.0	
		Machine Learning Engineer		MI	US
US	7	0.68	203479.0	193900.0	
		Data Analyst		SE	CA
CA	7	0.68	109198.0	130000.0	
		Data Engineer		MI	GR
GR	6	0.59	70920.0	73546.0	
		Data Engineer		SE	GB
GB	6	0.59	88682.0	89281.0	
		Data Engineer		MI	ES
ES	6	0.59	70920.0	73546.0	
		Data Scientist		SE	ES
ES	6	0.59	43460.0	43460.0	
		Data Engineer		EN	IN
IN	5	0.49	19215.0	17022.0	
		Data Scientist		MI	IN
IN	5	0.49	26236.0	30523.0	
		Data Scientist		MI	DE
DE	5	0.49	82179.0	88654.0	
		Data Scientist		EN	FR
FR	4	0.39	44382.0	44781.0	
		Data Scientist		EX	US
US	4	0.39	197188.0	192500.0	
		Data Scientist		EN	IN
IN	4	0.39	24712.0	25646.0	
		Data Scientist		MI	NL
NL	4	0.39	83265.0	81426.0	
		Data Analyst		EN	CA

CA	3	0.29	53221.0	52000.0	
Machine Learning Engineer				MI	GB
GB	3	0.29	112461.0	116976.0	
Data Analyst				MI	GR
GR	3	0.29	31182.0	31520.0	
Machine Learning Engineer				EN	GB
GB	3	0.29	40168.0	35093.0	
Data Engineer				EN	GB
GB	3	0.29	45913.0	45390.0	
Data Scientist				MI	ES
ES	3	0.29	41137.0	38776.0	
Data Engineer				SE	PR
PR	2	0.20	167500.0	167500.0	
Data Scientist				MI	CA
CA	2	0.20	71686.0	71686.0	
Data Analyst				MI	ES
ES	2	0.20	36773.0	36773.0	
Data Analyst				MI	CA
CA	2	0.20	80000.0	80000.0	
Data Analyst				EX	US
US	2	0.20	120000.0	120000.0	
Machine Learning Engineer				SE	IN
IN	2	0.20	45304.0	45304.0	
Machine Learning Engineer				SE	PR
PR	2	0.20	167500.0	167500.0	
Data Scientist				EN	BE
BE	2	0.20	68030.0	68030.0	
Data Scientist				EN	CA
CA	2	0.20	51417.0	51417.0	
Data Scientist				EN	DE
DE	2	0.20	55997.0	55997.0	
Data Analyst				EN	FR
FR	2	0.20	43735.0	43735.0	
Data Engineer				SE	CA
CA	2	0.20	161600.0	161600.0	
Data Engineer				MI	TR
TR	2	0.20	20060.0	20060.0	
Data Engineer				MI	DE
DE	2	0.20	62484.0	62484.0	
Data Analyst				SE	CF
CF	2	0.20	48609.0	48609.0	
Data Analyst				SE	GB
GB	2	0.20	73880.0	73880.0	
Data Analyst				SE	ES
ES	2	0.20	43602.0	43602.0	
Machine Learning Engineer				SE	AE
AE	2	0.20	92500.0	92500.0	
Data Scientist				SE	FR
FR	2	0.20	65438.0	65438.0	
Data Scientist				SE	IE
IE	2	0.20	142500.0	142500.0	

		Data Engineer		EN	PK
DE	2	0.20	55108.0	55108.0	
		Data Engineer		MI	FR
FR	2	0.20	67640.0	67640.0	
Machine Learning Engineer				SE	CA
CA	2	0.20	105000.0	105000.0	
		Data Scientist		SE	GB
GB	2	0.20	104663.0	104663.0	
		Data Scientist		SE	TR
TR	1	0.10	20171.0	20171.0	
		Data Scientist		MI	TR
TR	1	0.10	25000.0	25000.0	
		Data Scientist		SE	BR
US	1	0.10	100000.0	100000.0	
Machine Learning Engineer				EN	CH
CH	1	0.10	56536.0	56536.0	
		Data Scientist		SE	ES
GB	1	0.10	88256.0	88256.0	
		Data Scientist		SE	GR
US	1	0.10	68428.0	68428.0	
		Data Scientist		SE	AT
AT	1	0.10	91237.0	91237.0	
Machine Learning Engineer				MI	FR
DE	1	0.10	84053.0	84053.0	
Machine Learning Engineer				EN	CO
CO	1	0.10	21844.0	21844.0	
Machine Learning Engineer				EN	DE
DE	1	0.10	24823.0	24823.0	
Machine Learning Engineer				SE	PT
US	1	0.10	150000.0	150000.0	
Machine Learning Engineer				SE	NL
NL	1	0.10	59888.0	59888.0	
Machine Learning Engineer				SE	IE
IE	1	0.10	68293.0	68293.0	
Machine Learning Engineer				SE	HR
HR	1	0.10	45618.0	45618.0	
Machine Learning Engineer				SE	FI
FI	1	0.10	63040.0	63040.0	
Machine Learning Engineer				SE	DE
DE	1	0.10	94564.0	94564.0	
Machine Learning Engineer				SE	BE
BE	1	0.10	82744.0	82744.0	
Machine Learning Engineer				MI	SI
SI	1	0.10	24823.0	24823.0	
Machine Learning Engineer				MI	PL
PL	1	0.10	46597.0	46597.0	
Machine Learning Engineer				MI	NL
NL	1	0.10	96578.0	96578.0	
Machine Learning Engineer				MI	JP
JP	1	0.10	74000.0	74000.0	
Machine Learning Engineer				MI	IT

IT	1	0.10	51064.0	51064.0	
Machine Learning Engineer				MI	IN
IN	1	0.10	20984.0	20984.0	
Data Scientist				MI	RU
US	1	0.10	48000.0	48000.0	
Machine Learning Engineer				MI	ES
ES	1	0.10	47282.0	47282.0	
Machine Learning Engineer				MI	BE
BE	1	0.10	88654.0	88654.0	
Machine Learning Engineer				MI	AU
AU	1	0.10	83864.0	83864.0	
Machine Learning Engineer				EN	NL
DE	1	0.10	85000.0	85000.0	
Machine Learning Engineer				EN	IN
IN	1	0.10	20000.0	20000.0	
Data Scientist				MI	SG
IL	1	0.10	119059.0	119059.0	
Data Analyst				EN	AR
AR	1	0.10	50000.0	50000.0	
Data Scientist				MI	RS
DE	1	0.10	25532.0	25532.0	
Data Analyst				SE	DE
DE	1	0.10	63831.0	63831.0	
Data Engineer				MI	NL
NL	1	0.10	45391.0	45391.0	
Data Engineer				MI	MT
MT	1	0.10	28369.0	28369.0	
Data Engineer				MI	HK
GB	1	0.10	66022.0	66022.0	
Data Engineer				MI	ES
US	1	0.10	130800.0	130800.0	
Data Engineer				MI	AT
AT	1	0.10	74130.0	74130.0	
Data Engineer				EN	NL
NL	1	0.10	59888.0	59888.0	
Data Engineer				EN	JP
JP	1	0.10	41689.0	41689.0	
Data Engineer				EN	DE
DE	1	0.10	65013.0	65013.0	
Data Analyst				SE	PH
PH	1	0.10	50000.0	50000.0	
Data Analyst				SE	BG
US	1	0.10	80000.0	80000.0	
Data Engineer				MI	RO
US	1	0.10	26005.0	26005.0	
Data Analyst				MI	SG
SG	1	0.10	65257.0	65257.0	
Data Analyst				MI	PK
PK	1	0.10	8000.0	8000.0	
Data Analyst				MI	IN
IN	1	0.10	5723.0	5723.0	

		Data Analyst		MI	FR
FR	1	0.10	46759.0	46759.0	
		Data Analyst		EN	PT
PT	1	0.10	22809.0	22809.0	
		Data Analyst		EN	NG
NG	1	0.10	10000.0	10000.0	
		Data Analyst		EN	IN
IN	1	0.10	6072.0	6072.0	
		Data Analyst		EN	ID
ID	1	0.10	15000.0	15000.0	
		Data Analyst		EN	FR
IN	1	0.10	6359.0	6359.0	
		Data Engineer		MI	PL
PL	1	0.10	28476.0	28476.0	
		Data Engineer		SE	ES
US	1	0.10	193000.0	193000.0	
		Data Scientist		MI	PL
PL	1	0.10	33609.0	33609.0	
		Data Scientist		MI	CL
CL	1	0.10	40038.0	40038.0	
		Data Scientist		MI	PH
US	1	0.10	45760.0	45760.0	
		Data Scientist		MI	NG
NG	1	0.10	50000.0	50000.0	
		Data Scientist		MI	IN
US	1	0.10	5679.0	5679.0	
		Data Scientist		MI	IN
ID	1	0.10	53416.0	53416.0	
		Data Scientist		MI	HU
HU	1	0.10	35735.0	35735.0	
		Data Scientist		MI	HK
HK	1	0.10	65062.0	65062.0	
		Data Scientist		MI	FR
LU	1	0.10	62726.0	62726.0	
		Data Scientist		MI	FR
FR	1	0.10	42197.0	42197.0	
		Data Scientist		MI	DE
AT	1	0.10	61467.0	61467.0	
		Data Scientist		MI	CH
CH	1	0.10	120402.0	120402.0	
		Data Engineer		SE	GR
GR	1	0.10	47899.0	47899.0	
		Data Scientist		MI	BR
BR	1	0.10	12901.0	12901.0	
		Data Analyst		EN	BR
BR	1	0.10	7799.0	7799.0	
		Data Scientist		EN	UA
UA	1	0.10	13400.0	13400.0	
		Data Scientist		EN	JP
MY	1	0.10	40000.0	40000.0	
		Data Scientist		EN	ES



ES	1	0.10	31520.0	31520.0	
		Data Scientist		EN	AU
AU	1	0.10	83171.0	83171.0	
		Data Engineer		EX	ES
ES	1	0.10	79833.0	79833.0	
		Data Engineer		SE	RO
GB	1	0.10	76833.0	76833.0	
		Data Engineer		SE	MX
MX	1	0.10	33511.0	33511.0	
		Data Scientist		EN	US
DE	1	0.10	50000.0	50000.0	

```
/var/folders/jh/z981c7zj0vz0gmyfc8mhdxd0000gn/T/ipykernel_11047/346195264
2.py:5: FutureWarning: The default of observed=False is deprecated and
will be changed to True in a future version of pandas. Pass observed=False
to retain current behavior or observed=True to adopt the future default
and silence this warning.
```

```
salaries_data_frame.groupby(
```

```
# 1. Create new samples with STRING values and updated realistic salary
fields
```

```
new_samples = pd.DataFrame(
    [
        {
            "work_year": 2023,
            "experience_level": "SE",
            "employment_type": "FT",
            "job_title": "Data Scientist",
            "employee_residence": "US", ## United States
            "company_location": "US", ## United States
            "company_size": "M",
            "remote_ratio": 100,
            "salary": 160000,
            "salary_currency": "USD",
            "salary_in_usd": 160000,
        },
        {
            "work_year": 2023,
            "experience_level": "MI",
            "employment_type": "FT",
            "job_title": "Data Engineer",
            "employee_residence": "GB", ## United Kingdom
            "company_location": "GB", ## United Kingdom
            "company_size": "L",
            "remote_ratio": 50,
            "salary": 82528.0,
            "salary_currency": "USD",
            "salary_in_usd": 82528.0,
        },
    ],
```

```

        {
            "work_year": 2023,
            "experience_level": "EN",
            "employment_type": "FT",
            "job_title": "Data Analyst",
            "employee_residence": "BR", ## Brazil
            "company_location": "BR", ## Brazil
            "company_size": "S",
            "remote_ratio": 0,
            "salary": 8000,
            "salary_currency": "USD",
            "salary_in_usd": 8000,
        },
    ]
)

# 2. Drop target columns (keep only features the model expects)
X_new = new_samples.drop(columns=["salary", "salary_currency",
"salary_in_usd"])

# 3. Apply the SAME preprocessing pipeline you used for training
X_new_processed = preprocessor_2.transform(X_new)

# 4. Predict using the trained model
predictions = models_2["LinearRegression"].predict(X_new_processed)

# 5. Attach predictions back
new_samples["predicted_salary_usd"] = predictions.round(2)
import numpy as np

# 6. Compute error percentage
new_samples["error_percentage"] = (
    (new_samples["predicted_salary_usd"] - new_samples["salary_in_usd"])
    / new_samples["salary_in_usd"]
    * 100
).round(2)

# 7. Add comment about acceptability (e.g., <20% is okay in salary
prediction)
new_samples["comment"] = np.where(
    new_samples["error_percentage"].abs() <= 20, "✅ Acceptable", "⚠ High
Error"
)

# 8. Print results clearly
print(new_samples.to_string(index=False, line_width=10000))

```

```

work_year experience_level employment_type      job_title
employee_residence company_location company_size remote_ratio salary
salary_currency salary_in_usd predicted_salary_usd error_percentage
comment

```

2023	SE	FT Data Scientist	
US	US	M	100 160000.0
160000.0	161299.83	0.81	✓ Acceptable
2023	MI	FT Data Engineer	
GB	GB	L	50 82528.0
82528.0	78889.00	-4.41	✓ Acceptable
2023	EN	FT Data Analyst	
BR	BR	S	0 8000.0
8000.0	-23555.56	-394.44	△ High Error